



Improving I/O Performances of HPC Application Gysela

Alexandre Lou-Poueyou, Francieli Boito, Meline Trochon, Luan Teylo

**TECHNICAL
REPORT**

N° XXXX

MONTH YEAR

Project-Teams TADaaM



Improving I/O Performances of HPC Application Gysela

Alexandre Lou-Poueyou*, Francieli Boito[†], Meline Trochon[‡],

Luan Teylo [§]

Project-Teams TADaaM

Technical Report n° XXXX — MONTH YEAR — 6 pages

Abstract: In modern High-Performance Computing (HPC) systems, Input/Output (I/O) performance is a major bottleneck for large-scale scientific simulations. As computational power and data volumes scale rapidly, optimizing I/O workflows becomes critical to efficiently leverage supercomputing resources. Scientific applications heavily rely on Parallel File Systems (PFS) to access vast datasets concurrently. However, achieved performance is highly sensitive to PFS configurations, high-level I/O strategies, and application-specific access patterns. Understanding and tuning the interactions between these layers is crucial for maximizing global simulation throughput. This work focuses on optimizing the checkpointing performance of Gysela, a 5D gyrokinetic simulation code developed by the CEA for nuclear fusion research. Gysela generates massive datasets and relies on periodic checkpointing for fault tolerance and job restartability. Using `gys_io`, a lightweight mini-app replicating Gysela's parallel I/O patterns, we systematically evaluate and compare different parallel I/O strategies—namely Single Shared File (SSF), Subfilng, and File Per Process (FPP)—alongside key PFS configuration parameters. Our experiments, conducted on divers clusters, analyze the performance trade-offs of these strategies and parameters under various scaling conditions. The results identify optimal parallel configurations, clarify the impact of PFS striping on each I/O strategy, and provide actionable optimization recommendations to significantly reduce checkpointing overhead in future Gysela production runs.

Key-words: HPC, Parallel I/O, Performances Tuning, Parallel File Systems

* Computer Science Master's Degree Student at University of Bordeaux - Intern at TADaaM team

[†] Associate Professor at University of Bordeaux & permanent researcher at TADaaM team

[‡] PhD student at TADaaM team

[§] Permanent researcher at TADaaM team

RESEARCH CENTRE
Centre Inria de l'Université de Bordeaux

200 avenue de la Vieille Tour
33405 Talence Cedex

Contents

1	Introduction	4
2	Context	6
3	Background	6
4	Experimental Set-up	6

1 Introduction

In modern High-Performance Computing (HPC) systems, Input/Output (I/O) performance has become nearly critical as computational power itself. As supercomputer advance toward exascale capabilities, computational throughput has increased exponentially, yet storage bandwidth has lagged behind. This growing imbalance creates a fundamental bottleneck: while modern processors achieve teraflop- scale performance, parallel file systems typically deliver only tens to hundreds of gigabytes per second in aggregate bandwidth. Figure 1 quantifies this imbalance using data from major supercomputers from 2000 to 2025. The I/O bandwidth per unit of compute (measured in GBps per PetaFlop) has declined from ~ 300 GBps/PF on ASCI Red (2000) to only ~ 6 GBps/PF on Frontier (2023) a 50-fold degradation. This trend accelerated sharply after 2015, demonstrating that I/O has become the dominant bottleneck for modern exascale systems.





Figure 1: Compute-to-I/O performance gap in top 500 supercomputers (2000–2025). The ratio of computational throughput to storage bandwidth has grown by an order of indicating that I/O is increasingly the limiting factor.

Large-scale scientific simulations generate massive data volumes requiring efficient I/O management. On parallel file systems, thousands of processes must coordinate concurrent access to shared storage while contending for limited bandwidth. Gysela is a 5D gyrokinetic simulation code developed by the CEA (Commissariat à l'énergie Atomique) for nuclear fusion research. It models plasma turbulence in tokamaks by solving the Vlasov equation in a five-dimensional phase space: three spatial coordinates (minor radius and two toroidal angles) and two velocity-space coordinates (parallel velocity and magnetic moment). Understanding and controlling the turbulence that degrades plasma confinement is crucial for the success of next-generation fusion reactors like ITER.

- I/O is almost important as computation in HPC.
- As supercomputer advance, computational power increases faster than storage bandwidth, making I/O operation a bottleneck.



ratio between computation TFlops and I/O GB/s for top 500 supercomputer from 2010 to 2025.

- Large-Scale simulations generate unprecedented volumes of datas, much be stored and retrieved efficiently.
 - Example like dataset for AI training (ChatGPT), weather or any kind of simulation
- I/O in parallel file system is a complex problem due to concurrent access, contention, and data distribution.
- Gysela is a 5D gyrokinetic simulation code developed by CEA for nuclear fusion research (plasma turbulence in tokamaks).
- Run on large amount of nodes and generate massives datasets
 - Example of dataset size for Gysela
 - -> Ask Virginie about diversives size for Gysela prod
- Due to long execution time and risk of hardware failure, Gysela relies on periodic (every X [ask virginie] iterations) checkpointing to save its state and allow clean restart
- Writing checkpoint is synchronus and blocks the simulation progress, wasting expensive resources.
- On production, checkpoint write time suffer from high variability sometimes up to x3 slower than average.
- We must optimize both the PFS configuration and the I/O strategy
- We use gys_io, a lightweight mini-application that replicates Gysela's 5D grid checkpointing pattern.
 - Explain more how gys_io works
- We leverage PDI (Parallel Data Interface) to implement different I/O strategies: Single Shared File (SSF), Subfiling, and File Per Process (FPP) by modifying the configuration file in yaml, for flexible execution and neither code modification and compilation.
- We submits tests and evaluate the performances with IOPS, a benchmarking tool developped by TADaaM team at Inria Bordeaux.
- Our eval is conducted on different cluster with different PFS : Plafrim (with BeegFS) and Adastral (with Lustre). This allow us to test different PFS configuration and analyze the impact of striping on each I/O strategy.

- This work aims to identify optimal configurations and provide actionable recommendations to reduce checkpointing overhead in future Gysela production runs.

2 Context

Text for the second section. This is closely related to the text in section 1 on page 4.

3 Background

4 Experimental Set-up

The Inria logo is a red, stylized script font. It is positioned inside a white rectangular box with rounded corners, which has a subtle drop shadow.

RESEARCH CENTRE
Centre Inria de l'Université de Bordeaux

200 avenue de la Vieille Tour
33405 Talence Cedex

Publisher
Inria
Domaine de Voluceau - Rocquencourt
BP 105 - 78153 Le Chesnay Cedex
inria.fr

ISSN 0249-0803